

k-NN on Amazon Fine Food reviews Dataset

Exercise :

1. Download Amazon Fine Food Reviews dataset from Kaggle. You may have to create a Kaggle account to download data. (<https://www.kaggle.com/snap/amazon-fine-food-reviews>)
2. Perform featurization, BoW, tf-idf, Avg Word2Vec, tf-idf-Word2Vec.
3. Split data into train and test using time based slicing as 70% train & 30% test.
4. Perform 10-fold cross validation to find optimal k.
5. Report test accuracy for all four featurization.
6. Write your observations in English as crisply and unambiguously as possible. Always quantify your results.

Information regarding data set :

1. **Title:** Amazon Fine Food Reviews Data
2. **Sources:** Stanford Network Analysis Project(SNAP)
3. **Relevant Information:** This dataset consists of reviews of fine foods from amazon. The data span a period of more than 10 years, including all ~568,454 reviews up to October 2012(Oct 1999 - Oct 2012). Reviews include product and user information, ratings, and a plain text review.
4. **Attribute Information:**
 - ProductId** - unique identifier for the product
 - UserId** - unique identifier for the user
 - ProfileName** - name of the user
 - HelpfulnessNumerator** - number of users who found the review helpful
 - HelpfulnessDenominator** - number of users who indicated whether they found the review helpful or not
 - Score** - rating between 1 and 5.(rating of 4 or 5 could be considered a positive review. A review of 1 or 2 could be considered negative. A review of 3 is neutral and ignored)
 - Time** - timestamp for the review
 - Summary** - brief summary of the review
 - Text** - text of the review

Objective :

It is a 2-class classification task, where we have to analyze, transform(BoW,TF-IDF,AVG-W2V,TF-IDF-W2V) and calculate probabilistic class label values, which evaluates whether a review is positive or negative.

```
In [2]: import warnings
warnings.filterwarnings("ignore", category=UserWarning,module='gensim')
warnings.filterwarnings("ignore", category=UserWarning)

from sklearn.exceptions import DataConversionWarning
warnings.filterwarnings(action='ignore', category=DataConversionWarning)

with warnings.catch_warnings():
    warnings.simplefilter("ignore")

import traceback
import sqlite3
import datetime as dt
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap
import seaborn as sn
import itertools
from tqdm import tqdm
from sklearn import preprocessing
from sklearn.preprocessing import StandardScaler
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import cross_validate
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.decomposition import TruncatedSVD
from gensim.models import Word2Vec
from prettytable import PrettyTable
from imblearn.over_sampling import SMOTE
from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score
from sklearn.metrics import make_scorer, accuracy_score, confusion_matrix, classification_report
```

(1) Load dataset :

```
In [4]: # Load 'finalDataSet.sqlite' in panda's dataframe.
# This dataset is already gone through data deduplication and text preprocessing, so it is approx ~364
K

# Create connection object to load sqlite dataset
connection = sqlite3.connect('finalDataSet.sqlite')

# Load data into pandas dataframe.
reviews_df = pd.read_sql_query(""" SELECT * FROM Reviews """,connection)

# Drop index column
reviews_df = reviews_df.drop(columns=['index'])

# Convert timestamp to datetime.
reviews_df['Time'] = reviews_df[['Time']].applymap(lambda x: dt.datetime.fromtimestamp(x))

# Sort the data on the basis of time.
reviews_df = reviews_df.sort_values(by=['Time'])

# Take first 100K sample of reviews
reviews_df = reviews_df.head(100000)

print("Dataset Shape : \n",reviews_df.shape)
print("\nColumn Names: \n",reviews_df.columns)
print("\nTarget Class label : ")
print(reviews_df['Score'].value_counts())
print()


# Split data into 70% training and 30% testing.
x_train_original,x_test_original,y_train_original,y_test_original = train_test_split(reviews_df['Clean
edText'].values,
reviews_df['Score'].values,
test_size=0.3,
shuffle=False,
random_state=0)

# Plot review counts
plot_count_values(reviews_df)
```

Dataset Shape :
(100000, 11)

Column Names:
Index(['Id', 'ProductId', 'UserId', 'ProfileName', 'HelpfulnessNumerator',
 'HelpfulnessDenominator', 'Score', 'Time', 'Summary', 'Text',
 'CleanedText'],
 dtype='object')

Target Class label :
positive 88009
negative 11991
Name: Score, dtype: int64



```

In [8]: ###--- All utility variables and functions ---###

# List of odd numbers from 0 to 30
neighbors = list(filter(lambda x: x % 2 != 0, list(range(0,30))))

# Training Error
train_error = []

# Test Error
test_error = []

# Test Error
list_k = []

# http://scikit-learn.org/stable/modules/model_evaluation.html#scoring-parameter
# for list allowed scoring values
scoring = {'acc': 'accuracy',
           'prec_macro': 'precision_macro',
           'rec_micro': 'recall_macro'}

scoring_parameter = "Accuracy, Precision, Recall"

# Target Classes
target_classes = ["negative", "positive"]

def get_optimal_k(x_train, y_train, algorithm_name):
    '''
    This function, plots error and k values and, then returns optimal k value.
    '''

    scores = dict()

    # Pretty table instance
    ptable = PrettyTable()
    ptable.title = "Optimal K : 10-Fold Cross Validation"
    ptable.field_names = ["K Value", "Cross Validation Scoring Mean", "Scoring Parameter Used"]

    # Perform 10-fold cross validation
    for k in neighbors:
        knn_classifier = KNeighborsClassifier(n_neighbors=k, algorithm=algorithm_name, n_jobs=-1)
        result = cross_val_score(knn_classifier, x_train, y_train, cv=10, scoring=custom_scorer)
        scores[k] = result.mean()
        ptable.add_row([k, scores[k], scoring_parameter])

    # Print pretty table values
    print(ptable)

    # Plot the value of alpha's(x-axis) and cross validation scoring(accuracy,precision,recall)(y-axis)
    plt.plot(scores.keys(), scores.values())
    plt.xlabel("Value of k for k-NN")
    plt.ylabel("Cross validated scoring - accuracy,precision,recall")
    plt.show()

    optimal_k = max(scores, key=scores.get)
    list_k.append(optimal_k)
    print("\nOptimal value of hyperparameter k is ", optimal_k)

    return optimal_k

def getScores(estimator, x, y):
    yPred = estimator.predict(x)
    return (accuracy_score(y, yPred),
            precision_score(y, yPred, pos_label=3, average='macro'),
            recall_score(y, yPred, pos_label=3, average='macro'))

def custom_scorer(estimator, x, y):
    a, p, r = getScores(estimator, x, y)
    return a+p+r

def apply_k_nn(algorithm_name, optimal_k, x_train, y_train):
    '''
    This function tries to fit the model and returns the corresponding classifier.
    '''

    # instantiate learning model k with optimal_k and specified algorithm(brute,kd-tree)
    knn_classifier = KNeighborsClassifier(n_neighbors=optimal_k, algorithm=algorithm_name, n_jobs=-1)

    # Fitting the model.
    knn_classifier.fit(x_train, y_train)

    return knn_classifier

```

```

def plot_count_values(reviews_df):
    sn.catplot(x="Score",kind='count',data=reviews_df,height=5)
    plt.title("Score Count", fontsize=18)
    plt.show()

def plot_confusion_matrix(cm, classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.tight_layout()
    plt.show()

def generate_report(optimal_k, y_test, predicted_y_test):
    """
    This function generate reports like recall,precision,f1-score,confusion matrix.
    """

    print()
    # Pretty table instance
    ptable = PrettyTable()
    ptable.title = "Classification Report with k = {0}".format(optimal_k)
    ptable.field_names = ["Class Label/Averages", "Precision", "Recall", "F1-Score", "Support"]
    report_dict = classification_report(y_test, predicted_y_test,output_dict = True)
    for key , value in report_dict.items():
        inner_dict = value
        ptable.add_row([key,inner_dict['precision'],inner_dict['recall'],inner_dict['f1-score'],inner_
dict['support']])

    # Print pretty table values
    print(ptable)

    print()
    print("\nAccuracy Score: {0}%".format(accuracy_score(y_test, predicted_y_test)*100))
    test_error.append(1-accuracy_score(y_test, predicted_y_test))
    print()
    cnf_mat = confusion_matrix(y_test, predicted_y_test)
    plt.figure()
    plot_confusion_matrix(cnf_mat, classes=target_classes,title='Confusion Matrix')
    TN = cnf_mat[0,0]
    FP = cnf_mat[0,1]
    FN = cnf_mat[1,0]
    TP = cnf_mat[1,1]

    # Sensitivity, hit rate, recall, or true positive rate
    TPR = TP/(TP+FN)

    # Specificity or true negative rate
    TNR = TN/(TN+FP)

    # Fall out or false positive rate
    FPR = FP/(FP+TN)

    # False negative rate
    FNR = FN/(TP+FN)

    # Overall accuracy
    ACC = (TP+TN)/(TP+FP+FN+TN)

```

```

print()
# Pretty table instance
ptable = PrettyTable()
ptable.title = "Confusion Matrix Report"
ptable.field_names = ['Term', 'Value']
ptable.add_row(["TP (True Positive)", TP])
ptable.add_row(["TN (True Negative)", TN])
ptable.add_row(["FP (False Positive)", FP])
ptable.add_row(["FN (False Negative)", FN])
ptable.add_row(["TPR (True Positive Rate)= TP/(TP+FN)", TPR])
ptable.add_row(["TNR (True Negative Rate)= TN/(TN+FP)", TNR])
ptable.add_row(["FPR (False Positive Rate)= FP/(FP+TN)", FPR])
ptable.add_row(["FNR (False Negative Rate)= FN/(TP+FN)", FNR])
ptable.add_row(["ACC (Accuracy)= (TP+TN)/(TP+FP+FN+TN)", ACC])

# Print pretty table values
print(ptable)

def conclude():
    ptable=PrettyTable()
    ptable.title = "****Conclusion****"
    ptable.field_names=["Vectorizer", "Model", "Algorithm", "Hyperparameter(k)", "Train Error", "Test Error"]
    ptable.add_row(["BoW",
                    "K-NN",
                    "Brute-Force Search",
                    list_k[0],
                    str(round(train_error[0], 2)*100)+"%",
                    str(round(test_error[0], 2)*100)+"%"])
    ptable.add_row(["BoW",
                    "K-NN",
                    "KD-Tree",
                    list_k[1],
                    str(round(train_error[1], 2)*100)+"%",
                    str(round(test_error[1], 2)*100)+"%"])

    ptable.add_row(["TF-IDF",
                    "K-NN",
                    "Brute-Force Search",
                    list_k[2],
                    str(round(train_error[2], 2)*100)+"%",
                    str(round(test_error[2], 2)*100)+"%"])
    ptable.add_row(["TF-IDF",
                    "K-NN",
                    "KD-Tree",
                    list_k[3],
                    str(round(train_error[3], 2)*100)+"%",
                    str(round(test_error[3], 2)*100)+"%"])

    ptable.add_row(["AVG W2V",
                    "K-NN",
                    "Brute-Force Search",
                    list_k[4],
                    str(round(train_error[4], 2)*100)+"%",
                    str(round(test_error[4], 2)*100)+"%"])
    ptable.add_row(["AVG W2V",
                    "K-NN",
                    "KD-Tree",
                    list_k[5],
                    str(round(train_error[5], 2)*100)+"%",
                    str(round(test_error[5], 2)*100)+"%"])

    ptable.add_row(["TF-IDF W2V",
                    "K-NN",
                    "Brute-Force Search",
                    list_k[6],
                    str(round(train_error[6], 2)*100)+"%",
                    str(round(test_error[6], 2)*100)+"%"])
    ptable.add_row(["TF-IDF W2V",
                    "K-NN",
                    "KD-Tree",
                    list_k[7],
                    str(round(train_error[7], 2)*100)+"%",
                    str(round(test_error[7], 2)*100)+"%"])

    print(ptable)

def run_knn(x_train,y_train,x_test,y_test,chosen_algorithm = None):
    algorithms = chosen_algorithm

    for algorithm_name in algorithms:

```

```

        if algorithm_name == 'brute':
            print()
            print("\n*****")
            print("----- Brute-Force Search Algorithm -----")
            print("*****\n")
        else:
            print("\n\n")
            print("\n*****")
            print("----- KD-Tree Algorithm -----")
            print("*****\n")

    if(algorithm_name == "kd_tree"):
        print("\nPerforming TruncatedSVD, which will return dense matrix with lower dimensions...\n")

        svd = TruncatedSVD(n_components = 100)
        x_train = svd.fit_transform(x_train)
        x_test = svd.transform(x_test)

    # Find optimal K
    optimal_k = get_optimal_k(x_train, y_train,algorithm_name)

    # Perform naive bayes
    classifier = apply_k_nn(algorithm_name,optimal_k,x_train,y_train)

    # Make class predictions for x_test
    # Also make class predictions for x_train(training error)
    predicted_y_test = classifier.predict(x_test)
    predicted_y_train = classifier.predict(x_train)
    train_error.append(1 - accuracy_score(y_train, predicted_y_train))

    # Generate report
    generate_report(optimal_k,y_test,predicted_y_test)

```

(2) Convert review text to vector representation and perform k-NN on the corresponding vector :

(2.1) Bag of Words (BoW) :

Note: KD-Tree is computationally very costly for high dimensional dataset, so we will use truncatedSVD to reduce the dimensionality of data.

```

In [5]: %%time

# Instantiate CountVectorizer (vectorizer)
bow_count_vectorizer = CountVectorizer()

# learn the 'vocabulary' of the training data (occurs in-place)
bow_count_vectorizer.fit(x_train_original)

# Transform training and testing data(features) into a 'document-term matrix' or 'row-column matrix'
x_train_dtm = bow_count_vectorizer.transform(x_train_original)
x_test_dtm = bow_count_vectorizer.transform(x_test_original)

# Data Standardization
sc = StandardScaler(with_mean=False)
x_train_dtm = sc.fit_transform(x_train_dtm)
x_test_dtm = sc.transform(x_test_dtm)

print("\nthe type of count vectorizer ",type(x_train_dtm))
print("the shape of BOW vectorizer ",x_train_dtm.get_shape())
print("the number of unique words ", x_train_dtm.get_shape()[1])

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of BOW vectorizer (70000, 36118)
the number of unique words 36118
Wall time: 4.15 s

```

```

In [9]: %%time


# Perform k-NN on dataset.(BoW)
try:
    choice = ['brute','kd_tree']
    run_knn(x_train_dtm, y_train_original, x_test_dtm, y_test_original, chosen_algorithm = choice)

except Exception:
    traceback.print_exc()

```

 ----- Brute-Force Search Algorithm -----


Optimal K : 10-Fold Cross Validation		
K Value	Cross Validation Scoring Mean	Scoring Parameter Used
1	1.9758874015270744	Accuracy, Precision, Recall
3	2.0815822586119976	Accuracy, Precision, Recall
5	2.13608057674377	Accuracy, Precision, Recall
7	2.1807375724696216	Accuracy, Precision, Recall
9	2.230342253354761	Accuracy, Precision, Recall
11	2.2487415924844205	Accuracy, Precision, Recall
13	2.2645068680048923	Accuracy, Precision, Recall
15	2.286919194114564	Accuracy, Precision, Recall
17	2.286403758688782	Accuracy, Precision, Recall
19	2.2722423587922207	Accuracy, Precision, Recall
21	2.288929131323467	Accuracy, Precision, Recall
23	2.247173970300085	Accuracy, Precision, Recall
25	2.2902042972637617	Accuracy, Precision, Recall
27	2.270967431881206	Accuracy, Precision, Recall
29	2.2840983517044955	Accuracy, Precision, Recall



Optimal value of hyperparameter k is 25

Classification Report with k = 25				
Class Lable/Averages	Precision	Recall	F1-Score	Support
negative	0.8666666666666667	0.006542526421741319	0.012987012987012986	3974
positive	0.8682682682682683	0.9998463075386153	0.929423530252161	26026
micro avg	0.8682666666666666	0.8682666666666666	0.8682666666666666	30000
macro avg	0.8674674674674675	0.5031944169801783	0.471205271619587	30000
weighted avg	0.8680561094427761	0.8682666666666666	0.8080262395984377	30000

Accuracy Score: 86.82666666666667%




Confusion Matrix Report	
Term	Value
TP (True Positive)	26022
TN (True Negative)	26
FP (False Positive)	3948
FN (False Negative)	4
TPR (True Positive Rate)= TP/(TP+FN)	0.9998463075386153
TNR (True Negative Rate)= TN/(TN+FP)	0.006542526421741319
FPR (False Positive Rate)= FP/(FP+TN)	0.9934574735782586
FNR (False Negative Rate)= FN/(TP+FN)	0.00015369246138476907
ACC (Accuracy)= (TP+TN)/(TP+FP+FN+TN)	0.8682666666666666

 ----- KD-Tree Algorithm -----

Performing TruncatedSVD, which will return dense matrix with lower dimensions...


Optimal K : 10-Fold Cross Validation		
K Value	Cross Validation Scoring Mean	Scoring Parameter Used
1	1.9420706220019224	Accuracy, Precision, Recall
3	1.9861669609343917	Accuracy, Precision, Recall
5	2.012114866530535	Accuracy, Precision, Recall
7	2.0345826150740725	Accuracy, Precision, Recall
9	2.0612883615796918	Accuracy, Precision, Recall
11	2.073735426431774	Accuracy, Precision, Recall
13	2.076114472030658	Accuracy, Precision, Recall
15	2.0584030803126963	Accuracy, Precision, Recall
17	2.042706242514274	Accuracy, Precision, Recall
19	2.0701151344159467	Accuracy, Precision, Recall
21	1.9956316762231854	Accuracy, Precision, Recall
23	2.0105681641733613	Accuracy, Precision, Recall
25	2.0146750432057488	Accuracy, Precision, Recall
27	2.0003888917585595	Accuracy, Precision, Recall
29	1.991116087445856	Accuracy, Precision, Recall



Optimal value of hyperparameter k is 13

Classification Report with k = 13				
Class Lable/Averages	Precision	Recall	F1-Score	Support
negative	0.5076923076923077	0.008303975842979365	0.016340678385739042	3974
positive	0.8683480875229664	0.9987704603089218	0.9290041278747699	26026
micro avg	0.8675666666666667	0.8675666666666667	0.8675666666666667	30000
macro avg	0.688020197607637	0.5035372180759505	0.4726724031302545	30000
weighted avg	0.8205732185547318	0.8675666666666667	0.8081066429324563	30000

Accuracy Score: 86.75666666666667%



Confusion Matrix Report	
Term	Value
TP (True Positive)	25994
TN (True Negative)	33
FP (False Positive)	3941
FN (False Negative)	32
TPR (True Positive Rate)= TP/(TP+FN)	0.9987704603089218
TNR (True Negative Rate)= TN/(TN+FP)	0.008303975842979365
FPR (False Positive Rate)= FP/(FP+TN)	0.9916960241570206
FNR (False Negative Rate)= FN/(TP+FN)	0.0012295396910781526
ACC (Accuracy)= (TP+TN)/(TP+FP+FN+TN)	0.8675666666666667

Wall time: 3h 57min 3s

(2.2) Term Frequency - Inverse Document Frequency (TF-IDF) :

```
In [10]: %%time

# Instantiate TfidfVectorizer (vectorizer)
tfidf_vectorizer = TfidfVectorizer(min_df = 3,ngram_range=(1,2))

# learn the 'vocabulary' of the training data (occurs in-place)
tfidf_vectorizer.fit(x_train_original)

# Transform training and testing data(features) into a 'document-term matrix' or 'row-column matrix'
x_train_dtm = tfidf_vectorizer.transform(x_train_original)
x_test_dtm = tfidf_vectorizer.transform(x_test_original)

# Data Standardization
x_train_dtm = sc.fit_transform(x_train_dtm)
x_test_dtm = sc.transform(x_test_dtm)

print("\nthe type of count vectorizer ",type(x_train_dtm))
print("the shape of TF-IDF vectorizer ",x_train_dtm.get_shape())
print("the number of unique words ", x_train_dtm.get_shape()[1])

the type of count vectorizer <class 'scipy.sparse.csr.csr_matrix'>
the shape of TF-IDF vectorizer (70000, 162272)
the number of unique words 162272
Wall time: 14 s
```

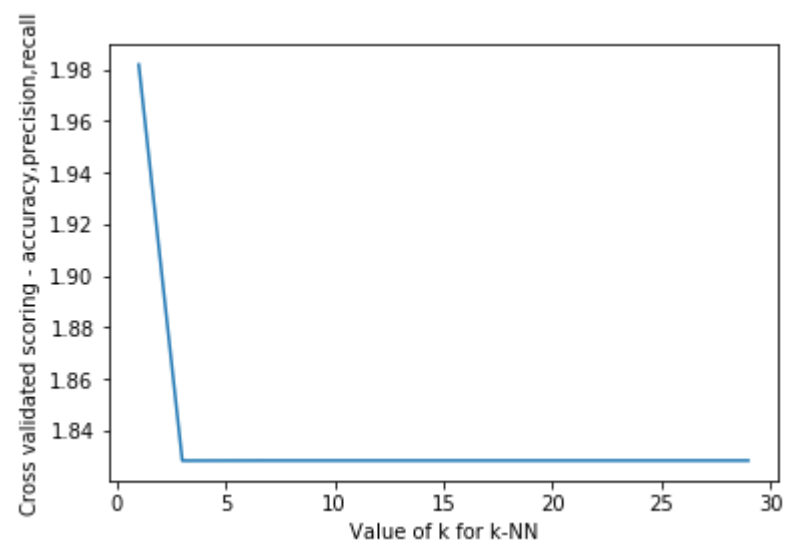
```
In [11]: %%time

# Perform k-NN on dataset.(BoW)
try:
    choice = ['brute','kd_tree']
    run_knn(x_train_dtm, y_train_original, x_test_dtm, y_test_original, chosen_algorithm = choice)

except Exception:
    traceback.print_exc()
```

 ----- Brute-Force Search Algorithm -----


Optimal K : 10-Fold Cross Validation		
K Value	Cross Validation Scoring Mean	Scoring Parameter Used
1	1.981832064764698	Accuracy, Precision, Recall
3	1.8281608045616953	Accuracy, Precision, Recall
5	1.8281839821964732	Accuracy, Precision, Recall
7	1.8282071499377555	Accuracy, Precision, Recall
9	1.8282071499377555	Accuracy, Precision, Recall
11	1.8282071499377555	Accuracy, Precision, Recall
13	1.8282071499377555	Accuracy, Precision, Recall
15	1.8282071499377555	Accuracy, Precision, Recall
17	1.8282071499377555	Accuracy, Precision, Recall
19	1.8282071499377555	Accuracy, Precision, Recall
21	1.8282071499377555	Accuracy, Precision, Recall
23	1.8282071499377555	Accuracy, Precision, Recall
25	1.8282071499377555	Accuracy, Precision, Recall
27	1.8282071499377555	Accuracy, Precision, Recall
29	1.8282071499377555	Accuracy, Precision, Recall



Optimal value of hyperparameter k is 1

Classification Report with k = 1					
Class Lable/Averages	Precision	Recall	F1-Score	Support	
negative	0.30714285714285716	0.010820332159033719	0.020904229460379193	3974	
positive	0.868352310783657	0.9962729578114193	0.9279247038614322	26026	
micro avg	0.8657333333333334	0.8657333333333334	0.8657333333333334	30000	
macro avg	0.5877475839632571	0.5035466449852265	0.4744144666609057	30000	
weighted avg	0.794010765158039	0.8657333333333334	0.8077747250191061	30000	

Accuracy Score: 86.57333333333334%




Confusion Matrix Report	
Term	Value
TP (True Positive)	25929
TN (True Negative)	43
FP (False Positive)	3931
FN (False Negative)	97
TPR (True Positive Rate)= TP/(TP+FN)	0.9962729578114193
TNR (True Negative Rate)= TN/(TN+FP)	0.010820332159033719
FPR (False Positive Rate)= FP/(FP+TN)	0.9891796678409662
FNR (False Negative Rate)= FN/(TP+FN)	0.00372704218858065
ACC (Accuracy)= (TP+TN)/(TP+FP+FN+TN)	0.8657333333333334

 ----- KD-Tree Algorithm -----

Performing TruncatedSVD, which will return dense matrix with lower dimensions...


Optimal K : 10-Fold Cross Validation		
K Value	Cross Validation Scoring Mean	Scoring Parameter Used
1	1.8881627977720357	Accuracy, Precision, Recall
3	1.9377068912015578	Accuracy, Precision, Recall
5	1.9590974181422816	Accuracy, Precision, Recall
7	1.987215725156068	Accuracy, Precision, Recall
9	2.027064477757251	Accuracy, Precision, Recall
11	2.0628837031265888	Accuracy, Precision, Recall
13	2.0777778311313977	Accuracy, Precision, Recall
15	2.096177663906706	Accuracy, Precision, Recall
17	2.0965963271193226	Accuracy, Precision, Recall
19	2.047977027808475	Accuracy, Precision, Recall
21	2.071558979608495	Accuracy, Precision, Recall
23	2.0836608024837537	Accuracy, Precision, Recall
25	2.102290957491914	Accuracy, Precision, Recall
27	2.0623372129711104	Accuracy, Precision, Recall
29	2.0913380485521094	Accuracy, Precision, Recall



Optimal value of hyperparameter k is 25

Classification Report with k = 25				
Class Label/Averages	Precision	Recall	F1-Score	Support
negative	0.5454545454545454	0.003019627579265224	0.006006006006006006	3974
positive	0.8678364133698045	0.999615768846538	0.9290764945361045	26026
micro avg	0.8676	0.8676	0.8676	30000
macro avg	0.706645479412175	0.5013176982129016	0.46754125027105525	30000
weighted avg	0.8251315619332966	0.8676	0.8068004238221507	30000

Accuracy Score: 86.76%



Confusion Matrix Report	
Term	Value
TP (True Positive)	26016
TN (True Negative)	12
FP (False Positive)	3962
FN (False Negative)	10
TPR (True Positive Rate)= TP/(TP+FN)	0.999615768846538
TNR (True Negative Rate)= TN/(TN+FP)	0.003019627579265224
FPR (False Positive Rate)= FP/(FP+TN)	0.9969803724207348
FNR (False Negative Rate)= FN/(TP+FN)	0.0003842311534619227
ACC (Accuracy)= (TP+TN)/(TP+FP+FN+TN)	0.8676

Wall time: 4h 40min 18s

(2.3) Average Word2Vec :

```

In [19]: %%time

# Make list of list from training data
list_of_sentences_in_train=[]
for sentence in x_train_original:
    list_of_sentences_in_train.append(sentence.split())

# Make list of list from testing data - this will be useful when vectorizing testing data.
list_of_sentences_in_test=[]
for sentence in x_test_original:
    list_of_sentences_in_test.append(sentence.split())

print("Shape of training data : ",x_train_original.shape)
print("Shape of testing data : ",x_test_original.shape)
print("Number of sentences present in training data : ",len(list_of_sentences_in_train))
print("Number of sentences present in testing data : ",len(list_of_sentences_in_test))

# Generate model.
w2v_model = Word2Vec(list_of_sentences_in_train,min_count=3,size=101, workers=6)

w2v_words = list(w2v_model.wv.vocab)
print("Length of vocabulary : ",len(w2v_words))

# Prepare train vectorizer using trained word2vec model
train_list = []
for sentence in tqdm(list_of_sentences_in_train,unit=" sentence",desc='Average Word2Vec - Train data'):
    word_2_vec = np.zeros(101)
    cnt_words = 0
    for word in sentence:
        if word in w2v_words:
            vec = w2v_model.wv[word]
            word_2_vec += vec
            cnt_words += 1
    if cnt_words != 0 :
        word_2_vec /= cnt_words
    train_list.append(word_2_vec)

# Prepare test vectorizer using trained word2vec model
test_list = []
for sentence in tqdm(list_of_sentences_in_test,unit=" sentence",desc='Average Word2Vec - Test data'):
    word_2_vec = np.zeros(101)
    cnt_words = 0
    for word in sentence:
        if word in w2v_words:
            vec = w2v_model.wv[word]
            word_2_vec += vec
            cnt_words += 1
    if cnt_words != 0 :
        word_2_vec /= cnt_words
    test_list.append(word_2_vec)

avg_w2v_train = np.array(train_list)
avg_w2v_test = np.array(test_list)

Shape of training data : (70000,)
Shape of testing data : (30000,)
Number of sentences present in training data : 70000
Number of sentences present in testing data : 30000
Length of vocabulary : 13792

Average Word2Vec - Train data: 100%|████████████████████████████████████████| 70000/70000 [02:02<00:00, 570.21 sentence/s]
Average Word2Vec - Test data: 100%|████████████████████████████████████████| 30000/30000 [00:57<00:00, 519.26 sentence/s]

Wall time: 3min 14s

```

```

In [20]: %%time


# Perform k-NN on dataset.(AVG-W2V)
try:
    choice = ['brute','kd_tree']
    run_knn(avg_w2v_train, y_train_original, avg_w2v_test, y_test_original, chosen_algorithm = choice)

except Exception:
    traceback.print_exc()

```

 ----- Brute-Force Search Algorithm -----


Optimal K : 10-Fold Cross Validation		
K Value	Cross Validation Scoring Mean	Scoring Parameter Used
1	2.1628714206952155	Accuracy, Precision, Recall
3	2.2551569323385814	Accuracy, Precision, Recall
5	2.2871062765765644	Accuracy, Precision, Recall
7	2.2996565698373597	Accuracy, Precision, Recall
9	2.308852176035892	Accuracy, Precision, Recall
11	2.31511878310089	Accuracy, Precision, Recall
13	2.31323855605532	Accuracy, Precision, Recall
15	2.3127818380065657	Accuracy, Precision, Recall
17	2.3108915845934193	Accuracy, Precision, Recall
19	2.3108114653081158	Accuracy, Precision, Recall
21	2.310647107732236	Accuracy, Precision, Recall
23	2.3146211169019155	Accuracy, Precision, Recall
25	2.312155879338108	Accuracy, Precision, Recall
27	2.314523813922711	Accuracy, Precision, Recall
29	2.312627095480022	Accuracy, Precision, Recall



Optimal value of hyperparameter k is 11

Classification Report with k = 11				
Class Label/Averages	Precision	Recall	F1-Score	Support
negative	0.7321848081440877	0.23527931555108203	0.35612264330603693	3974
positive	0.8941962886885074	0.9868592945516023	0.938245447405432	26026
micro avg	0.8873	0.8873	0.8873	30000
macro avg	0.8131905484162976	0.6110693050513422	0.6471840453557345	30000
weighted avg	0.8727351678990566	0.8873	0.8611335799557321	30000

Accuracy Score: 88.73%




Confusion Matrix Report	
Term	Value
TP (True Positive)	25684
TN (True Negative)	935
FP (False Positive)	3039
FN (False Negative)	342
TPR (True Positive Rate)= TP/(TP+FN)	0.9868592945516023
TNR (True Negative Rate)= TN/(TN+FP)	0.23527931555108203
FPR (False Positive Rate)= FP/(FP+TN)	0.7647206844489179
FNR (False Negative Rate)= FN/(TP+FN)	0.013140705448397755
ACC (Accuracy)= (TP+TN)/(TP+FP+FN+TN)	0.8873

 ----- KD-Tree Algorithm -----

Performing TruncatedSVD, which will return dense matrix with lower dimensions...


Optimal K : 10-Fold Cross Validation		
K Value	Cross Validation Scoring Mean	Scoring Parameter Used
1	2.1621815915863367	Accuracy, Precision, Recall
3	2.255385587565956	Accuracy, Precision, Recall
5	2.2872566920414	Accuracy, Precision, Recall
7	2.299530946644346	Accuracy, Precision, Recall
9	2.3088421670702077	Accuracy, Precision, Recall
11	2.3155276359259878	Accuracy, Precision, Recall
13	2.3129902025637676	Accuracy, Precision, Recall
15	2.31296312596434	Accuracy, Precision, Recall
17	2.311428540292225	Accuracy, Precision, Recall
19	2.310997376329459	Accuracy, Precision, Recall
21	2.3108412271331895	Accuracy, Precision, Recall
23	2.3139363533046655	Accuracy, Precision, Recall
25	2.3125187980817055	Accuracy, Precision, Recall
27	2.3139119089597338	Accuracy, Precision, Recall
29	2.3125070976070115	Accuracy, Precision, Recall



Optimal value of hyperparameter k is 11

Classification Report with k = 11				
Class Lable/Averages	Precision	Recall	F1-Score	Support
negative	0.73046875	0.23527931555108203	0.35591929958127144	3974
positive	0.8941852367688022	0.9867440252055637	0.9381872648230007	26026
micro avg	0.8872	0.8872	0.8872	30000
macro avg	0.8123269933844011	0.6110116703783228	0.6470532822021361	30000
weighted avg	0.8724982594881615	0.8872	0.8610561683606464	30000

Accuracy Score: 88.72%



Confusion Matrix Report	
Term	Value
TP (True Positive)	25681
TN (True Negative)	935
FP (False Positive)	3039
FN (False Negative)	345
TPR (True Positive Rate)= TP/(TP+FN)	0.9867440252055637
TNR (True Negative Rate)= TN/(TN+FP)	0.23527931555108203
FPR (False Positive Rate)= FP/(FP+TN)	0.7647206844489179
FNR (False Negative Rate)= FN/(TP+FN)	0.013255974794436333
ACC (Accuracy)= (TP+TN)/(TP+FP+FN+TN)	0.8872

Wall time: 4h 35min 47s

(2.4) Term Frequency - Inverse Document Frequency Weighted Word2Vec (TF-IDF-Word2Vec) :

```

In [21]: %%time

# Make list of list from training data.
sentences_in_train=[]
for sentence in x_train_original:
    sentences_in_train.append(sentence.split())

# Make list of list from testing data - this will be useful when vectorizing testing data.
sentences_in_test=[]
for sentence in x_test_original:
    sentences_in_test.append(sentence.split())

# Generate model
w2v_model = Word2Vec(sentences_in_train,min_count=3,size=101, workers=6)

# Instantiate TfidfVectorizer
tfidf_vectorizer = TfidfVectorizer(min_df = 3,ngram_range=(1,2))

# Tokenize and build vocab
tfidf_vectorizer.fit(x_train_original)

# Encode document
x_train_matrix = tfidf_vectorizer.transform(x_train_original)

# Get feature names
feature_names = tfidf_vectorizer.get_feature_names()

# Dictionary with word as a key, and the idf as a value
dict_word_idf = dict(zip(feature_names, list(tfidf_vectorizer.idf_)))

# Prepare train vectorizer using trained word2vec model
train_list = []
row = 0
for sentence in tqdm(sentences_in_train,unit=" sentence",desc='TF-IDF Weighted Word2Vec - Train data'):
    word_2_vec = np.zeros(101)
    weight_tfidf_sum = 0
    for word in sentence:
        try:
            vec = w2v_model.wv[word]
            # dict_word_idf[word] = idf value of word in whole corpus
            # sentence.count(word) = tf value of word in this review
            tfidf_value = dict_word_idf[word]*sentence.count(word)
            word_2_vec += (vec * tfidf_value)
            weight_tfidf_sum += tfidf_value
        except:
            pass
    if weight_tfidf_sum != 0:
        word_2_vec /= weight_tfidf_sum
    train_list.append(word_2_vec)
    row += 1

# Prepare test vectorizer using trained word2vec model
test_list = []
row = 0
for sentence in tqdm(sentences_in_test, unit=" sentence",desc='TF-IDF Weighted Word2Vec - Test data'):
    word_2_vec = np.zeros(101)
    weight_tfidf_sum = 0
    for word in sentence:
        try:
            vec = w2v_model.wv[word]
            # dict_word_idf[word] = idf value of word in whole corpus
            # sentence.count(word) = tf value of word in this review
            tfidf_value = dict_word_idf[word]*sentence.count(word)
            word_2_vec += (vec * tfidf_value)
            weight_tfidf_sum += tfidf_value
        except:
            pass
    if weight_tfidf_sum != 0:
        word_2_vec /= weight_tfidf_sum
    test_list.append(word_2_vec)
    row += 1

tfidf_w2v_train = np.array(train_list)
tfidf_w2v_test = np.array(test_list)

TF-IDF Weighted Word2Vec - Train data: 100%|████████████████████| 70000/70000 [00:18<00:00, 38
71.72 sentence/s]
TF-IDF Weighted Word2Vec - Test data: 100%|████████████████████| 30000/30000 [00:08<00:00, 37
39.67 sentence/s]

```


Wall time: 44.7 s

```
In [22]: %%time

# Perform k-NN on dataset.(TF-IDF W2V)
try:
    choice = ['brute','kd_tree']
    run_knn(tfidf_w2v_train,y_train_original,tfidf_w2v_test,y_test_original, chosen_algorithm = choice)
except Exception:
    traceback.print_exc()
```

----- Brute-Force Search Algorithm -----


Optimal K : 10-Fold Cross Validation			
K Value	Cross Validation Scoring Mean	Scoring Parameter Used	
1	2.112375770805098	Accuracy, Precision, Recall	
3	2.1989783480024974	Accuracy, Precision, Recall	
5	2.2394154812564855	Accuracy, Precision, Recall	
7	2.2502454973864006	Accuracy, Precision, Recall	
9	2.2564495932403306	Accuracy, Precision, Recall	
11	2.263034368995337	Accuracy, Precision, Recall	
13	2.2668180412481425	Accuracy, Precision, Recall	
15	2.2760606575805857	Accuracy, Precision, Recall	
17	2.2759919566887463	Accuracy, Precision, Recall	
19	2.276980929467527	Accuracy, Precision, Recall	
21	2.2819985092175235	Accuracy, Precision, Recall	
23	2.290588236821965	Accuracy, Precision, Recall	
25	2.286489548865189	Accuracy, Precision, Recall	
27	2.2848238099389278	Accuracy, Precision, Recall	
29	2.283615954029004	Accuracy, Precision, Recall	



Optimal value of hyperparameter k is 23

Classification Report with k = 23				
Class Lable/Averages	Precision	Recall	F1-Score	Support
negative	0.7218710493046776	0.1436839456467036	0.23966421825813225	3974
positive	0.8834948132424938	0.9915469146238377	0.934407531456504	26026
micro avg	0.8792333333333333	0.8792333333333333	0.8792333333333333	30000
macro avg	0.8026829312735857	0.5676154301352706	0.5870358748573181	30000
weighted avg	0.862085051979531	0.8792333333333333	0.8423772005681597	30000

Accuracy Score: 87.92333333333333%




Confusion Matrix Report	
Term	Value
TP (True Positive)	25806
TN (True Negative)	571
FP (False Positive)	3403
FN (False Negative)	220
TPR (True Positive Rate)= TP/(TP+FN)	0.9915469146238377
TNR (True Negative Rate)= TN/(TN+FP)	0.1436839456467036
FPR (False Positive Rate)= FP/(FP+TN)	0.8563160543532964
FNR (False Negative Rate)= FN/(TP+FN)	0.0084530853761623
ACC (Accuracy)= (TP+TN)/(TP+FP+FN+TN)	0.8792333333333333

 ----- KD-Tree Algorithm -----

Performing TruncatedSVD, which will return dense matrix with lower dimensions...


Optimal K : 10-Fold Cross Validation		
K Value	Cross Validation Scoring Mean	Scoring Parameter Used
1	2.112494686523518	Accuracy, Precision, Recall
3	2.199271415364962	Accuracy, Precision, Recall
5	2.2400450838170576	Accuracy, Precision, Recall
7	2.250817509916746	Accuracy, Precision, Recall
9	2.257107056729584	Accuracy, Precision, Recall
11	2.262266203335425	Accuracy, Precision, Recall
13	2.267592487623287	Accuracy, Precision, Recall
15	2.2754764941569228	Accuracy, Precision, Recall
17	2.2764637824090928	Accuracy, Precision, Recall
19	2.2776903004508813	Accuracy, Precision, Recall
21	2.2816199244234854	Accuracy, Precision, Recall
23	2.28995441132665	Accuracy, Precision, Recall
25	2.2853453828025754	Accuracy, Precision, Recall
27	2.2858217341891574	Accuracy, Precision, Recall
29	2.282489937144562	Accuracy, Precision, Recall



Optimal value of hyperparameter k is 23

Classification Report with k = 23					
Class Lable/Averages	Precision	Recall	F1-Score	Support	
negative	0.7222222222222222	0.14393558127830902	0.2400335711288292	3974	
positive	0.8835250616269515	0.9915469146238377	0.9344244487091284	26026	
micro avg	0.8792666666666666	0.8792666666666666	0.8792666666666666	30000	
macro avg	0.8028736419245869	0.5677412479510734	0.5872290099189789	30000	
weighted avg	0.8621578121671384	0.8792666666666666	0.8424408037923248	30000	

Accuracy Score: 87.92666666666666%



Confusion Matrix Report	
Term	Value
TP (True Positive)	25806
TN (True Negative)	572
FP (False Positive)	3402
FN (False Negative)	220
TPR (True Positive Rate)= TP/(TP+FN)	0.9915469146238377
TNR (True Negative Rate)= TN/(TN+FP)	0.14393558127830902
FPR (False Positive Rate)= FP/(FP+TN)	0.856064418721691
FNR (False Negative Rate)= FN/(TP+FN)	0.0084530853761623
ACC (Accuracy)= (TP+TN)/(TP+FP+FN+TN)	0.8792666666666666

Wall time: 2h 57min 47s

Conclusion :

In [23]: conclude()

Conclusion						
Vectorizer	Model	Algorithm	Hyperparameter(k)	Train Error	Test Error	
BoW	K-NN	Brute-Force Search	25	11.0%	13.0%	
BoW	K-NN	KD-Tree	13	11.0%	13.0%	
TF-IDF	K-NN	Brute-Force Search	1	0.0%	13.0%	
TF-IDF	K-NN	KD-Tree	25	11.0%	13.0%	
AVG W2V	K-NN	Brute-Force Search	11	9.0%	11.0%	
AVG W2V	K-NN	KD-Tree	11	9.0%	11.0%	
TF-IDF W2V	K-NN	Brute-Force Search	23	10.0%	12.0%	
TF-IDF W2V	K-NN	KD-Tree	23	10.0%	12.0%	

Observations :

1. After applying k-NN over BoW,TF_IDF,AVG-Word2Vec and TF_IDF-Word2Vec vectors, we can easily

compare testing accuracies for different factorization.

2. We have sensible confusion matrix, but still it can be improved with more data points with balanced positive and negative points.
3. Because of the dense nature of the matrix, it is computationally very costly to take more number of samples to train under kd-tree, .
4. Different classification techniques(Like logistic regression) can be applied to find perfect classifier, which can further increase the testing accuracy.